



University of Trento

Part 1: Images and Videos

Nicola Conci
nicola.conci@unitn.it



Definition

Computer vision (from Wikipedia)

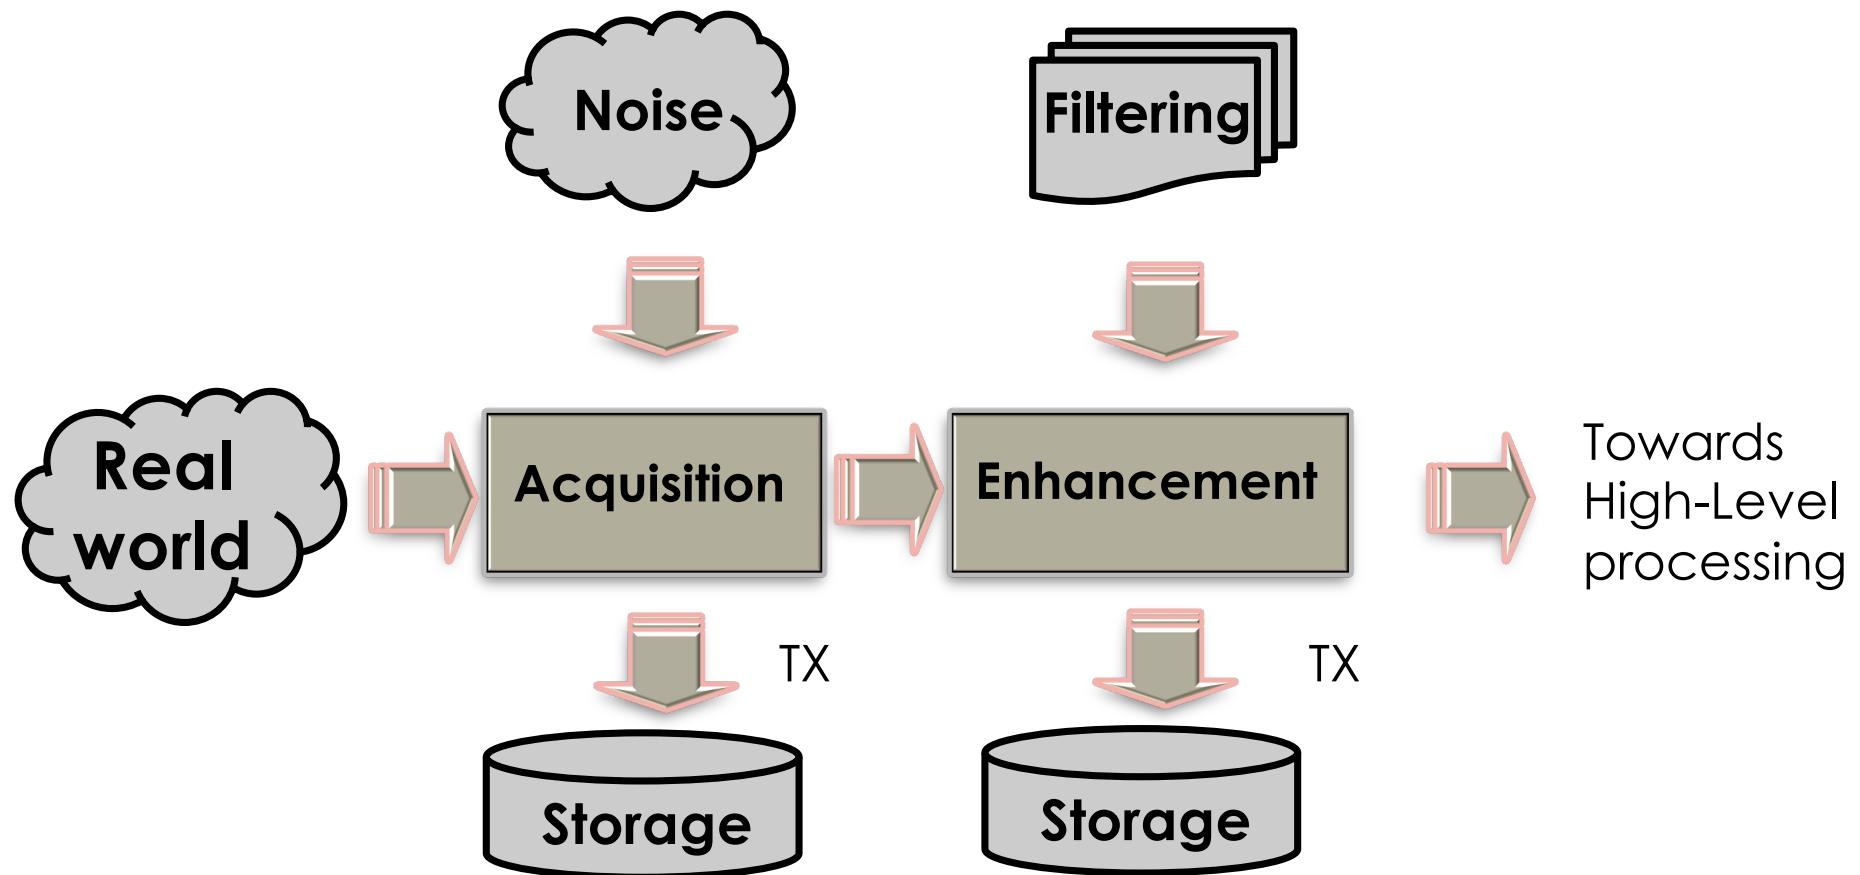
Computer vision is the **science and technology of machines that see**, where “see” in this case means that the machine is able to extract information from an image that is necessary to solve some task. [...] The image data can take many forms, such as video sequences, views from multiple cameras, or multi-dimensional data from a medical scanner. [...]

Examples of applications of computer vision include systems for:

- **Controlling processes** (e.g., an industrial robot or an autonomous vehicle)
- **Detecting events** (e.g. for visual surveillance or people counting)
- **Organizing information** (e.g. for indexing databases of images and image sequences)
- **Modeling objects or environments** (e.g., industrial inspection, medical image analysis or topographical modeling)
- **Interaction** (e.g., as the input to a device for computer-human interaction).



The processing chain



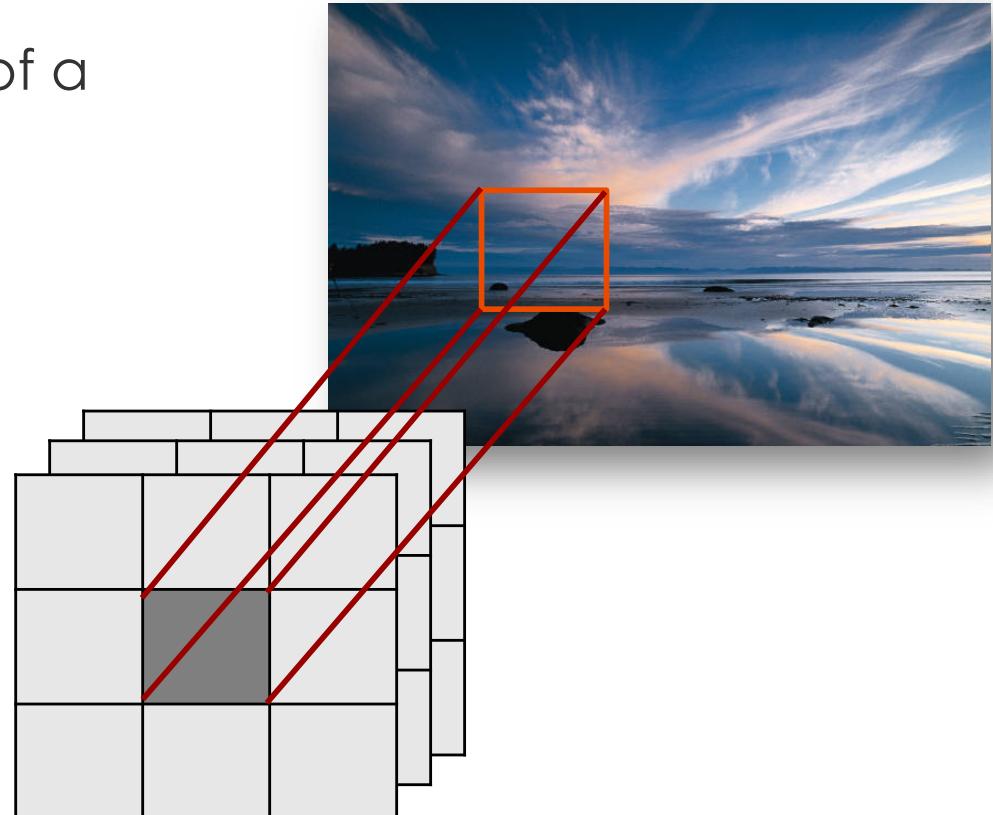


Acquisition

- Transformation of a physical signal into an electrical one (by means of a sensor)
 - **What** we measure: light intensity (B/W), wavelength (color), temperature (infrared), etc.
- It refers to the process of transferring a portion of the real 3D world onto a 2D surface
- It brings a continuous-parameter real world into a discrete-parameter one
- Representation in a standard format

Digital Images

- Collection of 2D coordinates
- Coordinates are known as pixels, picture elements
- A pixel represents a projection of a portion of the real world
- Pixels can be
 - Grayscale
 - One component, 8bit
 - Color
 - Typically 3 components, 24bits





Sampling

- The “real world” is a continuous function
- Analog video is a 1-D continuous function where one spatial dimension is mapped onto time by means of the scanning process
- Digital video is instead sampled in all three dimensions





Sampling in a nutshell

- Simplified case: 2D image
 - Continuous signal $s_c(x_1, x_2)$
 - Spatial rectangular sampling:
 - $x_1 = n_1 \Delta x_1$
 - $x_2 = n_2 \Delta x_2$
 - $s(n_1, n_2) = s_c(n_1 \Delta x_1, n_2 \Delta x_2)$
 - *or*

$$\begin{aligned}s_d(x_1, x_2) &= s_c(x_1, x_2) \sum_{n_1} \sum_{n_2} \delta(x_1 - n_1 \Delta x_1, x_2 - n_2 \Delta x_2) \\ &= \sum_{n_1} \sum_{n_2} s_c(n_1, n_2) \delta(x_1 - n_1 \Delta x_1, x_2 - n_2 \Delta x_2)\end{aligned}$$



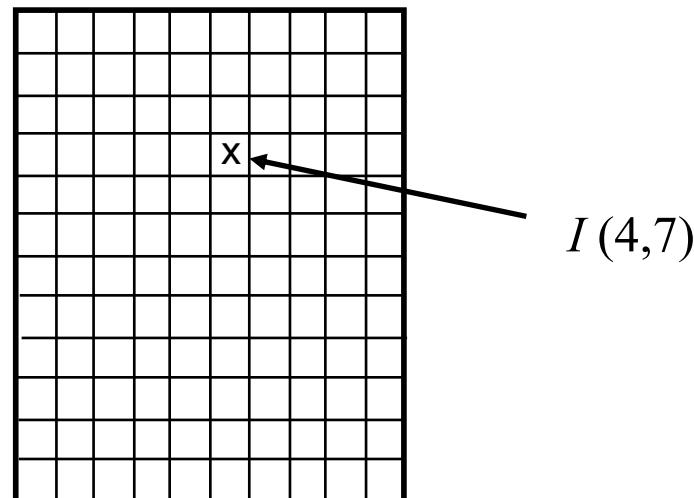
After sampling

- A 2D matrix of numbers, of pixels
- Once we have the digital format, we can manipulate data and:
 - Apply filters
 - Change colors
 - Store
 - Transmit



Handling images

- Pixels are numbered starting from the top left corner
- The value of a pixel in a certain position is defined as $I(r,c)$
 - r is the row index
 - c is the column index
- Numbering starts usually from 0 or 1, depending on habits...





Handling images

- Monochrome images
 - Values normalized in the range 0-1
 - 0 is black
 - 1 is white
 - The intensity is called *grey level*.
- Color pictures work the same way:
 - Three channels
 - Each of them in the range 0-1





Color

- What is color?
 - The attribute the human visual system associates to objects
 - A mathematical relationship that combines different wavelengths
- Why is it important?
 - To check whether something we see is what we expect
 - To recognize objects
 - To distinguish similar objects

Examples



- The bus is red, no doubt.
- What would the computer say?
 - Sure there's some red, BUT
 - It's also
 - Black
 - Yellow
 - White
 - Grey
 - ...

Examples



- (based on color) Can we say it's the same cat?



You would eat this:



And this?!

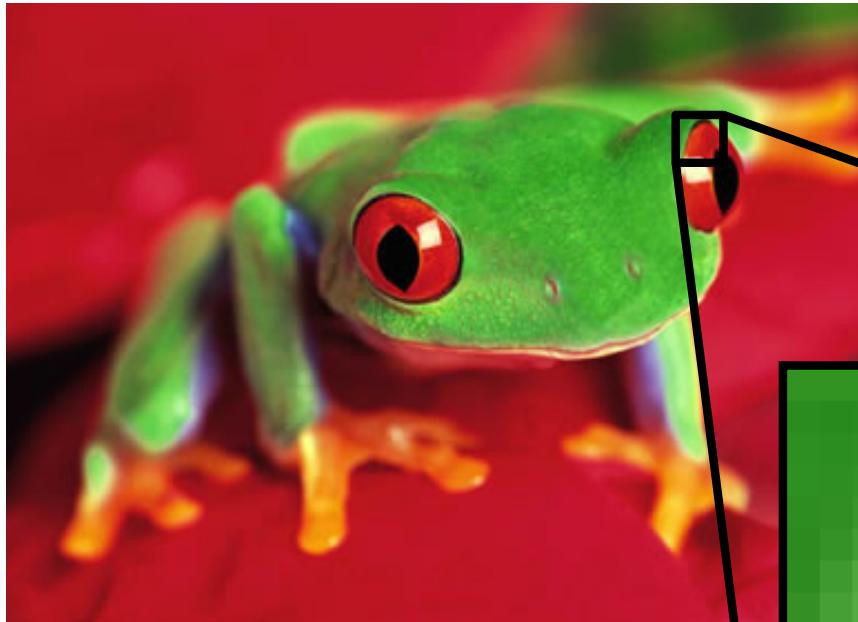


Color perception

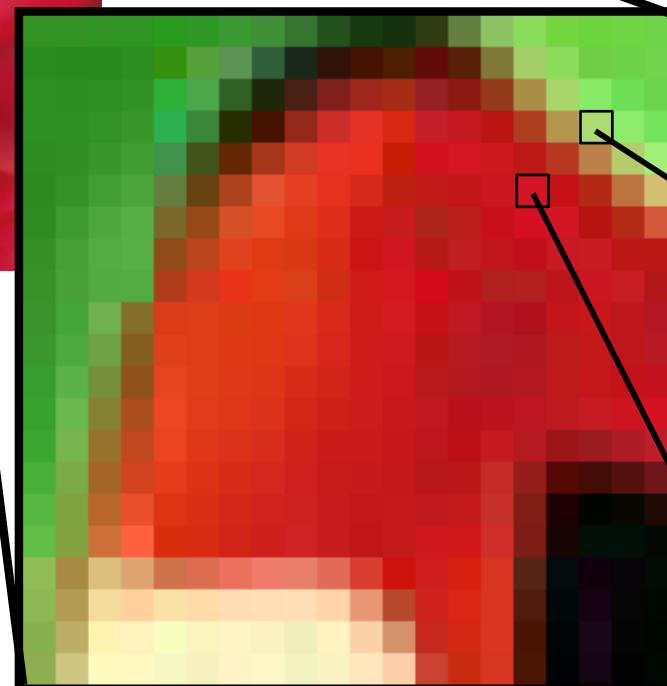
- The human eye “is a camera” with a focal length of about 20mm, where the iris controls the amount of light by adjusting the size of the pupil
- The perception of color is possible through cones in the fovea
- It has around 100M receptors
- Cones have peak responses on three main wavelengths
 - Red (700nm)
 - Green (546.1nm)
 - Blue (435.8nm)



Images: representation



RGB Picture
387x280 pixels
RGB, 8 bpp
Total: 2.600.640 bit



Detail, 20x20 pixel

**RGB
(136, 233, 102)**

**Each pixel is
an RGB triplet**

**RGB
(213, 21, 36)**



Image Compression

- Data can be:
 - Processed locally
 - Transmitted remotely
 - Archived on a storage unit
- Images and videos require a lot of bandwidth
- To compress a picture or a video, we need a codec
- A codec allows:
 - **CO**dec: encoding in the compressed domain
 - **coDEC**: decoding from the compressed domain
- Examples:
 - JPEG
 - MPEG
 - DIVX
- **Processing** is typically done in the **uncompressed** domain



Raw vs compressed

- Raw images
 - vector of pixels
 - Usually stored in 1D vector, of size $1 \times (N \times M)$
 - If 8 bit per pixel (bpp), storage requires $M \times N$ bytes
 - Color image in HDTV format needs: $1920 \times 1080 \times 3 > 6\text{MBytes}$
 - If it is a video at 25 pictures per second $\rightarrow 6 \times 25 = 150\text{MB}$
- Compressed images
 - Reduces the dimension of the file
 - With losses (most used)
 - Without losses (low compression)

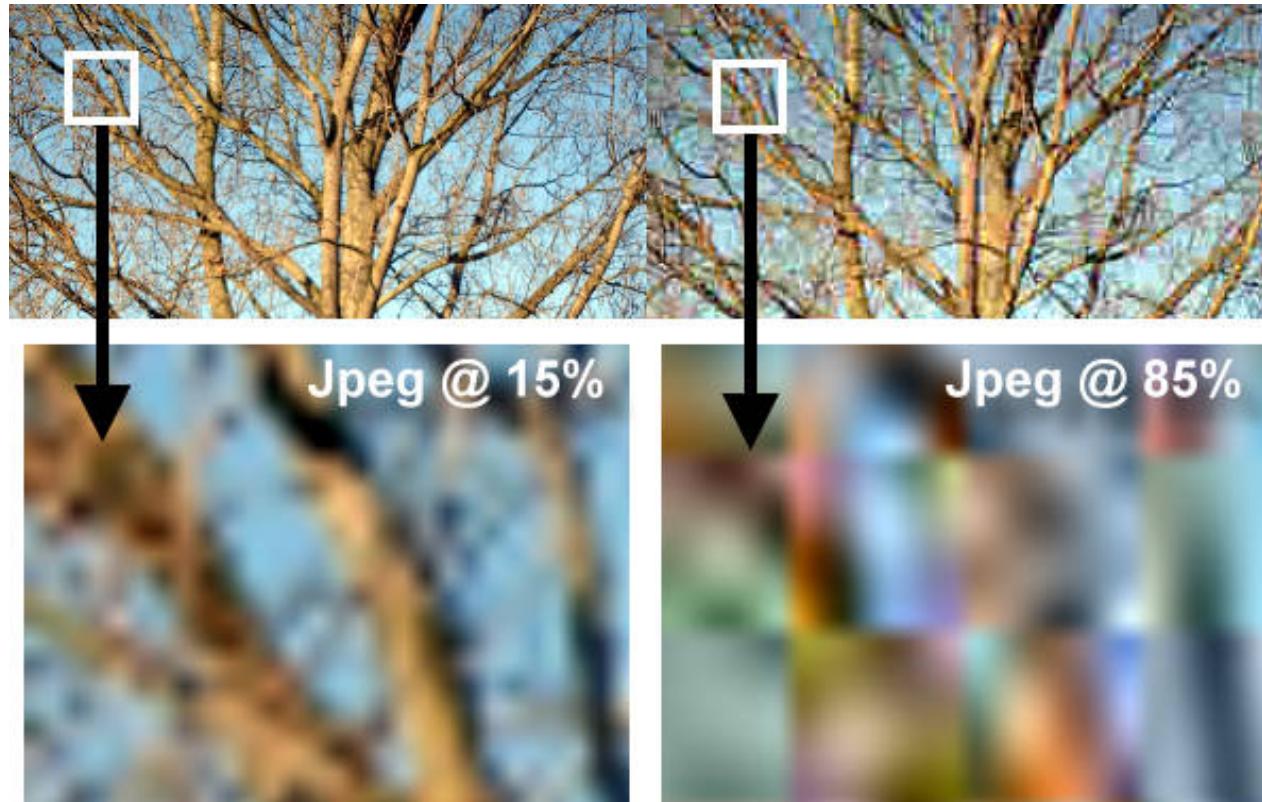


Compression

- Standards
 - JPEG, compression around 10-20 or more
 - 1 Mbyte can be stored in 50 Kbytes without losing too much in quality
 - The choice of the standard typically depends on the data
- Compression reduces quality
- Compression introduces visual artifacts
- Be careful when you compress data!

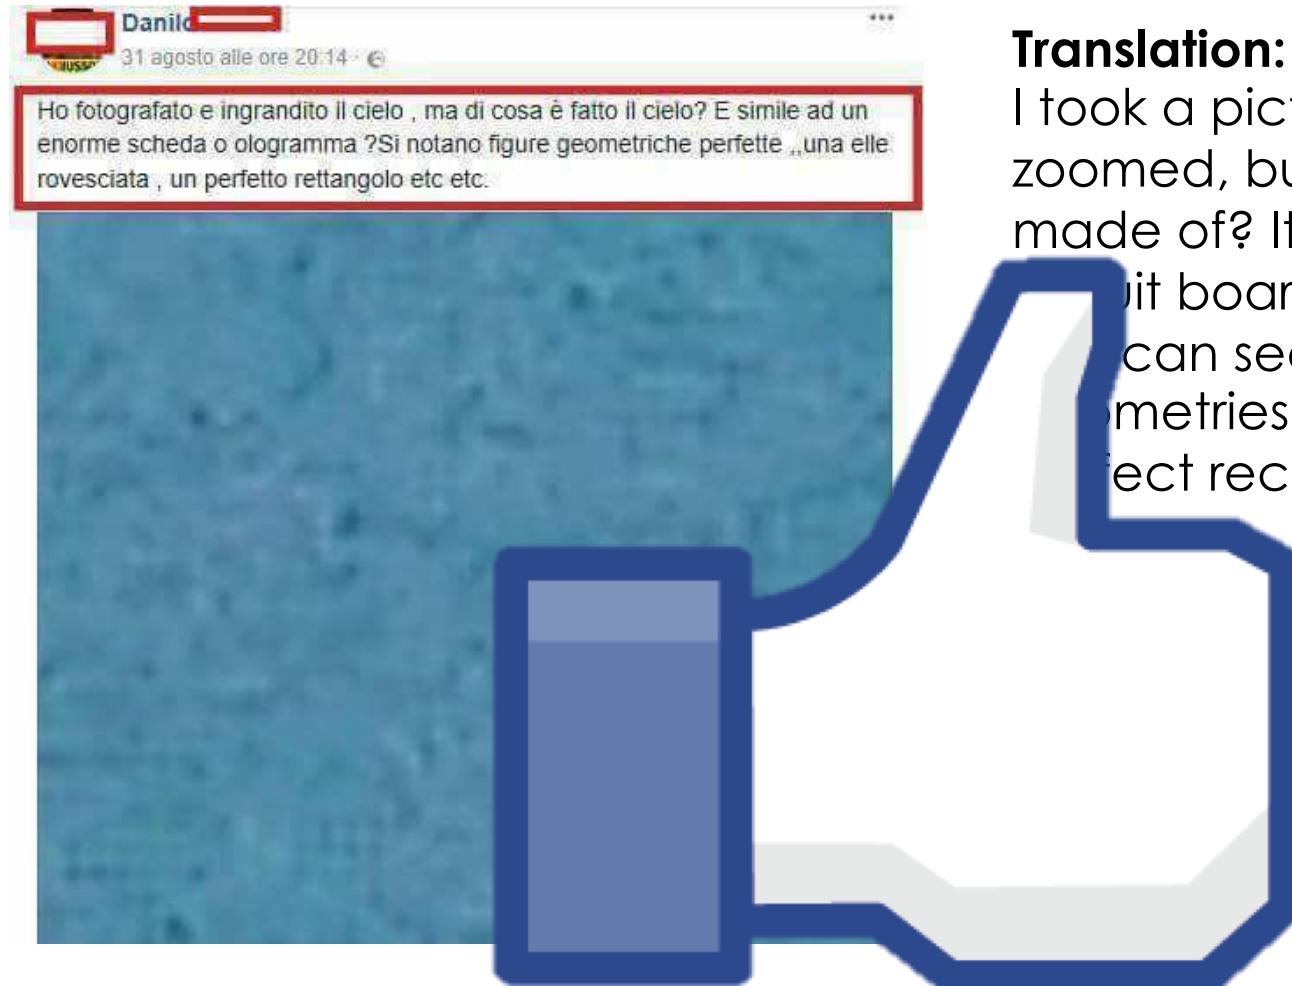


Compression: example



By increasing the compression ratio, artifacts appear such as blocking, blurring, chromatic aberrations

Compression or conspiracy?



Translation:

I took a picture of the sky and zoomed, but what is the sky made of? It's similar to a huge circuit board or a hologram? You can see perfect geometries, an L upside down, a perfect rectangle, etc etc.

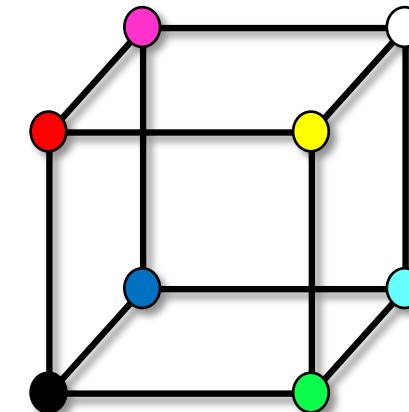
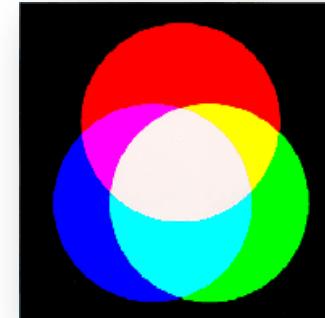
It compression
noise of the sensor



Additive Color Model

- Colored beams are projected onto a black surface
- Beams overlap
- Human eye receives the stimulus without generating interference
- The eye mixes the components and perceives the resulting color
- Starting from the primary colors RGB we can obtain:

- $R+G$ = Yellow
- $R+B$ = Magenta
- $B+G$ = Cyan
- $R+G+B$ = White



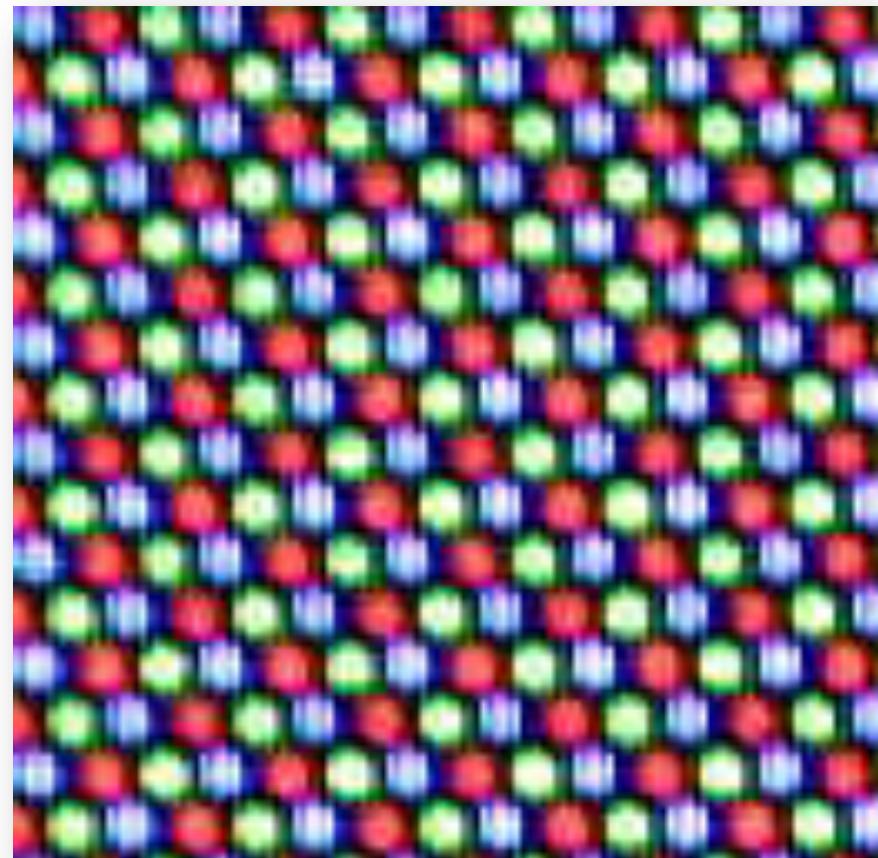
- **Subtractive color is the inverse process**



How are colors combined?

- Black: RGB (0,0,0)
- Green: RGB (0, 1, 0)
- Yellow: RGB (1, 1, 0)
- White: RGB (1, 1, 1)
- Grey: RGB (0.5, 0.5, 0.5)

Display





Color spaces: why RGB?

- Major response in the green component
- Red and Blue are *less* relevant
- Higher response to light than color
- Maybe a different representation would be more effective

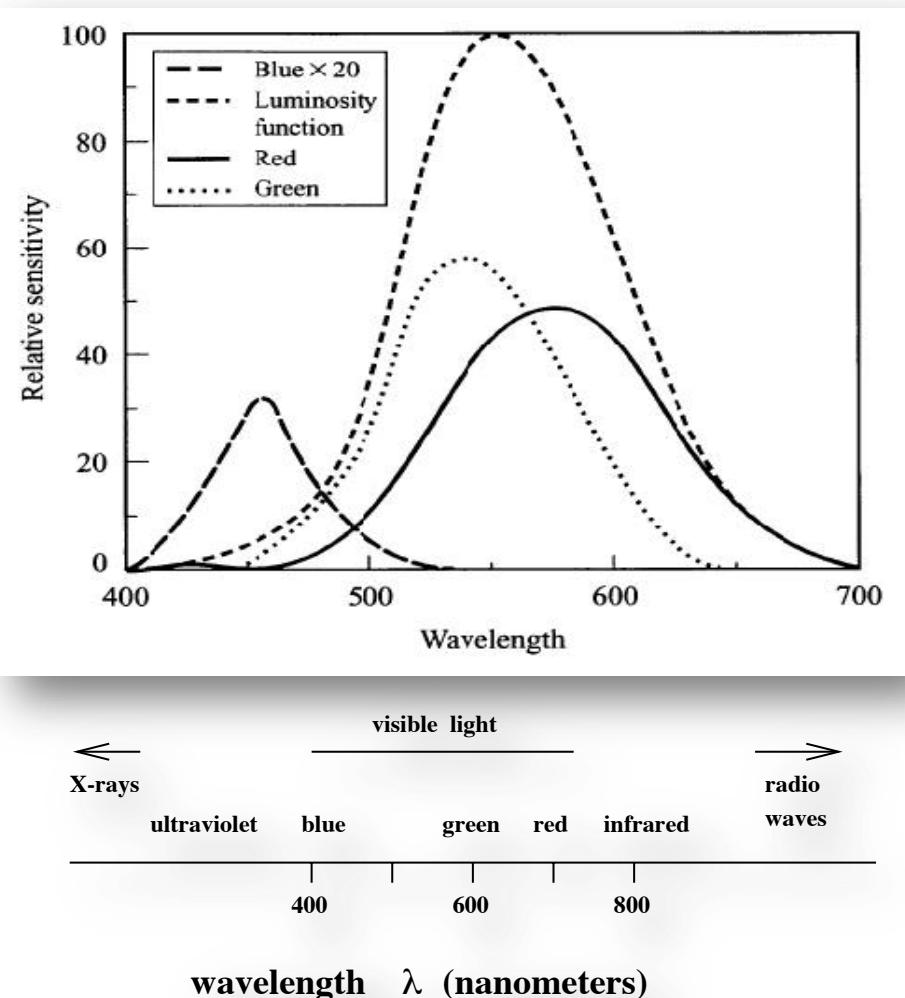




Image perception

- In general the human eye is more sensitive to luminance than color
- It is also more sensitive to **contrast variations** with respect to absolute values
- The internal square of the left image appears brighter than the one on the right
- Even though it's the same color...

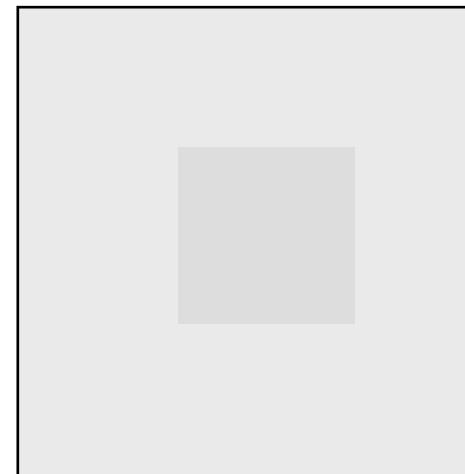
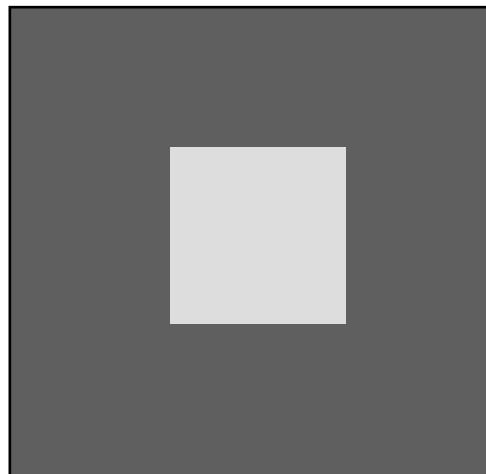




Image perception

- What is the luminance level of the stripes?
- Is it constant?
- It's actually constant, but perceived as over/undershoot of intensity because of the contrast





The Weber Law

- The minimum contrast variation that can be perceived btw two different luminance values f_s and f_o is given by the Weber law:

$$| f_s - f_o | / f_o = 0,02 \text{ (empirical number)}$$

- This means that in the range 0-1 the human eye is able to distinguish roughly 50 different variations.

Color spaces: RGB



- If we separate the three components and generate single images we notice:
 - Components are correlated
 - This means that the three greyscale images carry *almost* the same amount of information



Color spaces: YCbCr



- More effective, since luminance (Y) is separated from the chrominance components (Cb Cr)
- Compatible with the Human Visual System:
 - Rods (120M) are used to discriminate light levels
 - Cones (6-7M) respond to color stimuli in the RGB wavelengths.

Color spaces: YUV



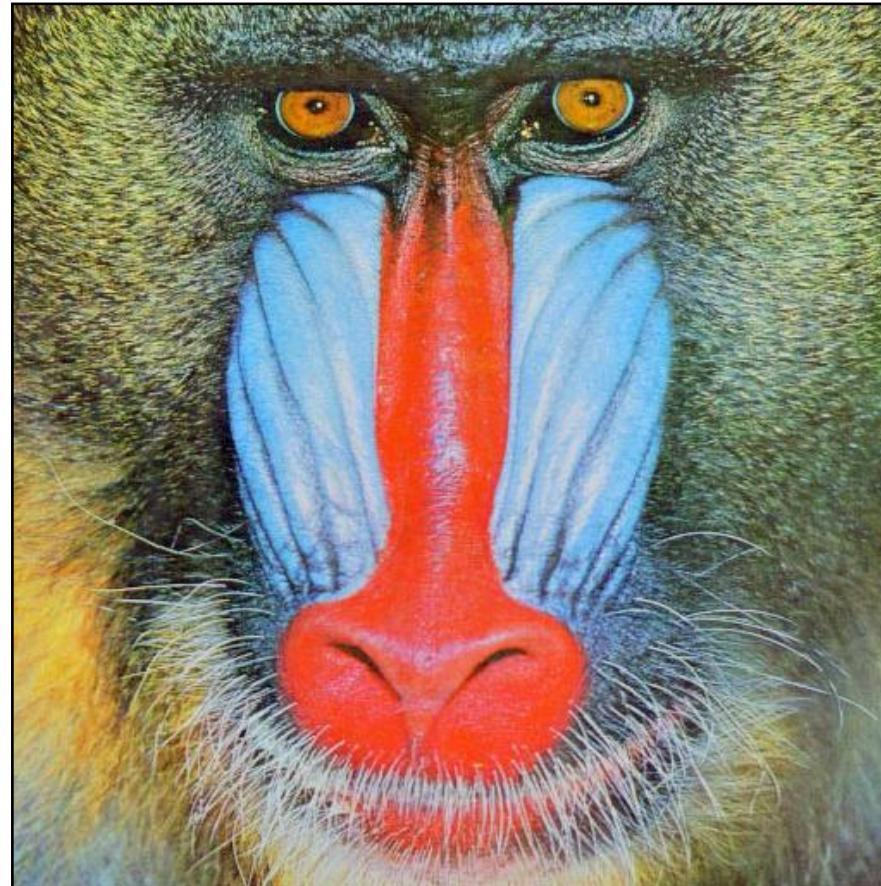
- Different formats:
 - 4:4:4 → each component is fully represented
 - 4:2:0 → chrominances are downsampled by a factor 2
 - ...
- YCbCr is a generalization of YUV, just a matter of conversion matrices
- Downsampling of color is done for convenience (storage)



Color representation



Original picture $\textbf{YC}_b\textbf{C}_r$, no downsampling



Color representation



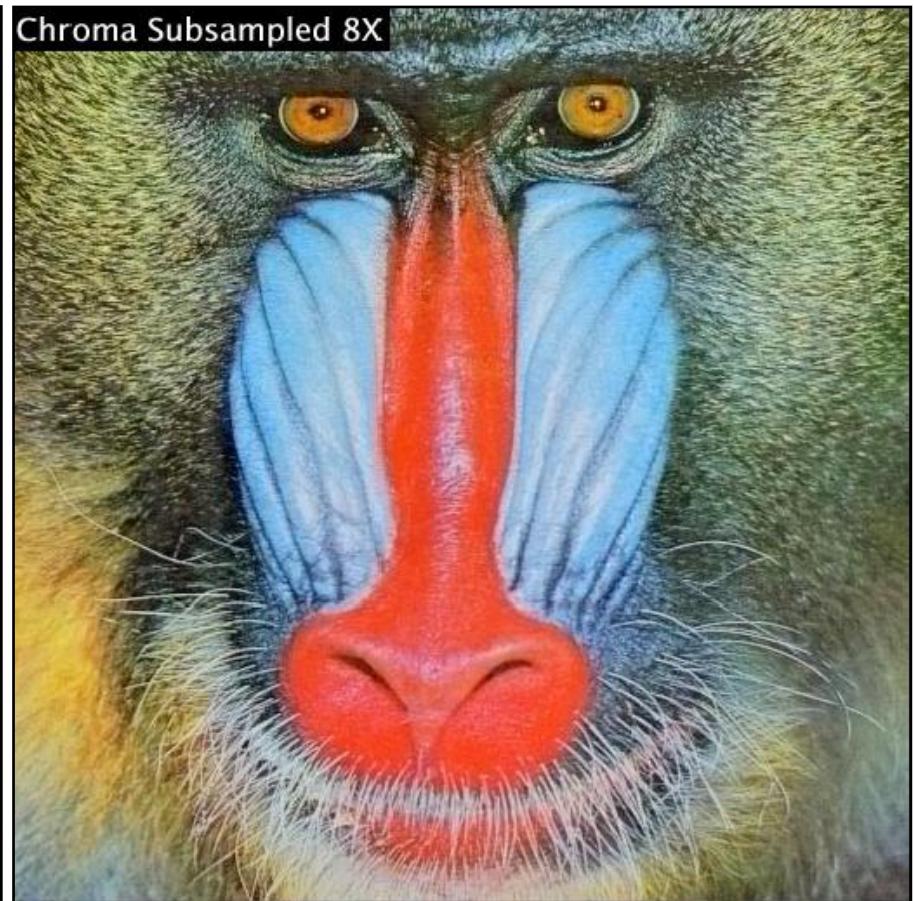
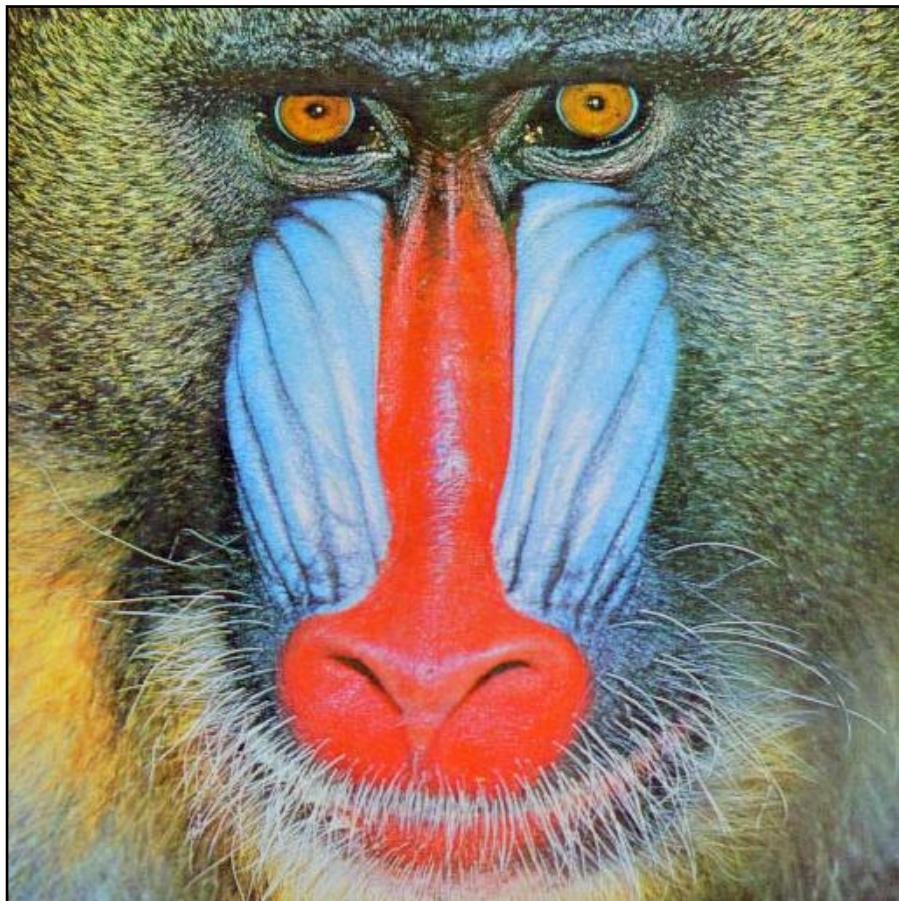
Image YC_bC_r downsampling on Y



Color representation



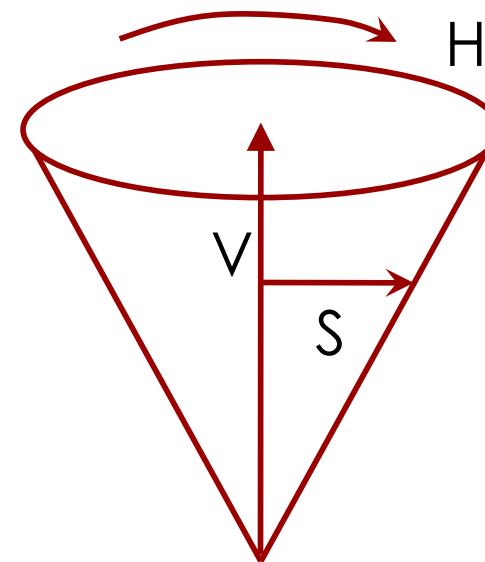
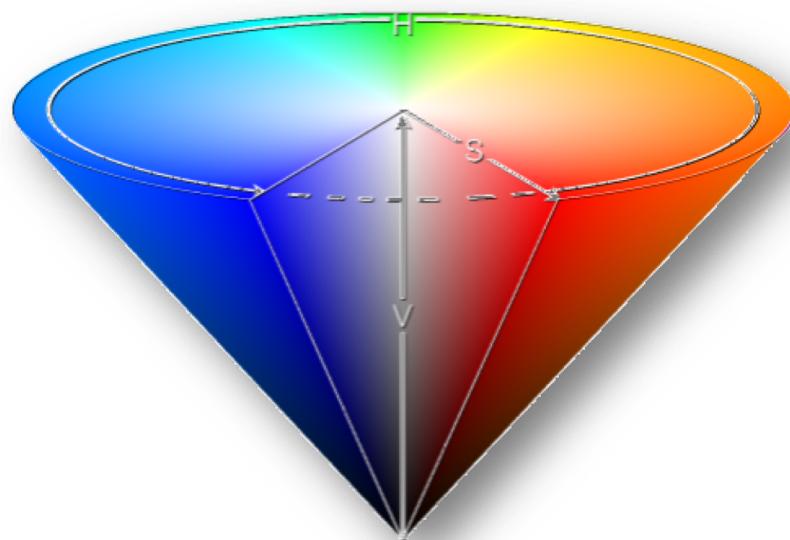
Image YC_bC_r , downsampling on C_bC_r





Color spaces: HSV

- Color is represented through:
 - Hue
 - Saturation
 - Value



Different color spaces for different applications



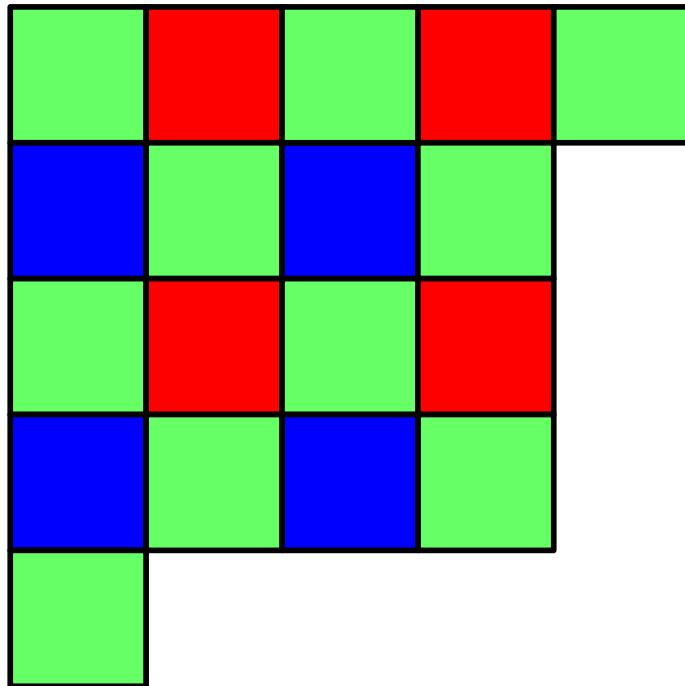
- RGB is used in general for visualization
 - In displays each pixel is composed by three phosphors (CRT) or LEDs (LCD)
- YUV is suitable for compression
 - We are less sensitive to chrominance variations
 - U and V can be downsampled
- HSV is robust for computer graphics and image analysis
 - H is the “color”
 - V is “brightness”
 - S is the “intensity”



Bayer Pattern

- In the acquisition phase, light is captured by the CCD (Charge Coupled Device)
- The CCD is an array of cells
- The best solution would be to have devices with 3 different CCDs
- Most cameras are single-chip
- To correctly exploit the human eye response:
 - Three types of photosensors
 - 50% green
 - 25% red
 - 25% blue

Bayer Pattern



- Green sensors are defined as *luminance-sensitive elements*
- Red and Blue sensors are defined as *chrominance-sensitive elements*



Quantization

- Like in the mono-dimensional case, signals need to be quantized
- Quantization implies the definition of a number of levels to define our signal
- Typically, the range 0-1 is quantized using 8bpp
- Other representations with 10-12 bpp are available



Why 8bpp?

- 8bpp represent 256 levels, which is fine for the human eye
- 7bpp (128 levels) would still be ok
- What happens if we quantize with less than 6bpp (64 levels, minimum to ensure “smooth” pictures)
- False contours appear → contouring

Contouring



- At 5bpp



Contouring



- At 4bpp





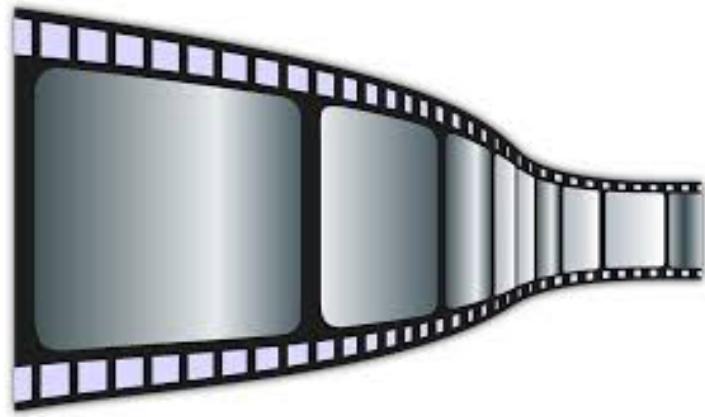
Limits of the 2D

- Still images provide a reliable information about static scenes
- We lose motion information:
 - Temporal evolution of the scene
 - Rapid changes
 - Dynamics of motion (qualitative and quantitative)
 - How subjects/objects relate one to each other
- Analyzing a video provides a more consistent representation of the scene
- It's closer to what humans do every day



What is a video?

- Sequence of 2D images that represents a projection of a moving 3D scene onto the video camera image plane.



It is expected that adjacent frames are strongly correlated



Resolution

- Display resolution: it determines the number of lines and rows in an image
- The term resolution could be misleading.
- It also refers to the actual resolution of the display in terms of *pixels per inch*. This is the pixel density, which is calculated by taking a known image and measuring the size (in inches) on the screen
- High density = high performance display = €€€



Resolution in images and video

- Images are up to 50MP
- Resolution in video is typically lower:
 - Full HD → $1920 \times 1080 = \sim 2\text{MP}$
 - 4k → $3840 \times 2160 = \sim 8.3\text{MP}$
 - 8k → $7680 \times 4320 = \sim 33\text{MP}$
- Reasons:
 - Videos can last hours
 - Videos have a higher frame rate, up to 60fps (also 120 or more)
 - Single frames last for a short time
 - Storage could be troublesome in terms of capacity and access to disk



Relevant features

- Once the image is acquired, what are the most relevant features we could be interested in?
 - Color and its distribution
 - Presence of edges and contours
 - ...
- If the analysis concerns a video, instead of an image, what is the added value?
 - Consistency of the features mentioned above over time
 - Evolution of the scene and objects displacement
 - Objects may enter/exit the scene

Examples

- Static background
- Low motion
- “Controlled” environment



Examples



- Noise
- Distortion
- Light artifacts



Examples

- Background with slow changes
- Consistent motion due to the presence of people
- Shadows
- Long temporal range of analysis





Comments

- Why are these scenes complex to analyze/process?
 - Edges are moving also if the scene “appears” static
 - Pixels are not constant even though they “look” constant
 - Shadows can be seen as moving objects depending on their intensity
 - Environmental conditions can significantly “disturb” the analysis modules
 - We have then highlights, reflections, occlusions, masking